

PA02 Assignment Report

Student Name: Dewansh Khandelwal

Roll Number: MT25067

Course: CSE638 - Graduate Systems

Assignment: Analysis of Network I/O Primitives using Perf

Submission Date: February 6, 2026

1. Assignment Completion Checklist

Component	Status	File(s)	Notes
Part A1: Two-Copy Implementation	✓ Complete	MT25067_PartA1_Server.c MT25067_PartA1_Client.c	Uses send()/recv() primitives
Part A2: One-Copy Implementation	✓ Complete	MT25067_PartA2_Server.c MT25067_PartA2_Client.c	Uses sendmsg() with iovec
Part A3: Zero-Copy Implementation	✓ Complete	MT25067_PartA3_Server.c MT25067_PartA3_Client.c	Uses sendmsg() with MSG_ZEROCOPY
Part B: Profiling & Measurement	✓ Complete	MT25067_ExperimentData.csv	48 experiments total
Part C: Automation Script	✓ Complete	MT25067_PartC_AutomationScript.sh	Fully automated, no manual intervention
Part D: Plotting	✓ Complete	MT25067_PartD_Plots.py	5 plots generated with hardcoded data
Part E: Analysis	✓ Complete	MT25067_PartE_Analysis.md	6 questions answered with technical depth
Makefile	✓ Complete	Makefile	Compiles all implementations
README	✓ Complete	README.md	Usage instructions and dependencies
GitHub Repository	✓ Complete	github.com/dewansh3255/GRS_PA02	Public repository

2. System Configuration

```

OS:           Ubuntu 22.04.5 LTS (Jammy Jellyfish)
Kernel:       6.8.0-90-generic
CPU:          Intel Core i7-12700 (12th Generation)
              - 12 cores (8 P-cores + 4 E-cores)
              - 20 threads with Hyperthreading
              - L3 Cache: 25 MB (shared)
Architecture: x86_64
Network:      Localhost loopback (lo interface)
Compiler:     GCC 11.4.0
Python:        3.10.12
Perf:         Linux perf_event subsystem

```

3. Implementation Overview

Part A1: Two-Copy Baseline

Copy Locations:

1. **User space:** Message fields → Serialization buffer (`serialize_message()`)
2. **Kernel space:** Serialization buffer → Kernel socket buffer (`send()` syscall)

Code Highlight:

```

// User-space copy (Copy #1)
serialize_message(msg, send_buffer, message_size);

// Kernel-space copy (Copy #2)
send(client_fd, send_buffer, message_size, 0);

```

Part A2: One-Copy Optimization

Eliminated Copy: User-space serialization removed using scatter-gather I/O.

Code Highlight:

```

// Setup iovec to point directly to message fields
struct iovec iov[8];
iov[0].iov_base = msg->field1; // No memcpy!
iov[0].iov_len = field_size;
// ... repeat for 8 fields

// Kernel gathers from scattered buffers (ONE copy total)
sendmsg(client_fd, &msghdr, 0);

```

Diagram:

```

A1: field1 → \
    field2 → → [buffer] → [kernel] → socket
    field8 → /
        (memcpy)      (memcpy)

A2: field1 → \
    field2 → → → → [kernel] → socket
    field8 → /           (gather & copy once)

```

Sample Execution Outputs**Single Experiment with Perf Profiling:**

```

Activities Terminal Feb 6 22:30 ⓘ
iiitd@iiitd-ThinkCentre-M70s-Gen-3: ~/Desktop/GR5_PA02$ sudo perf stat -e cycles,instructions,cache-misses,L1-dcache-load-misses,context-switches ./MT25067_PartA1_Server 4096 5000
Message size: 4096 bytes
Messages per client: 5000
Max clients: 1
Server listening on port 8080...
Client 1 connected
All 1 clients accepted. Waiting for transfers to complete...
[Thread 129910281184832] Handling client, sending 5000 messages of 4096 bytes
[Thread 129910281184832] Sent 20480000 bytes in 0.006 sec (29356.75 Mbps)
Performance counter stats for './MT25067_PartA1_Server 4096 5000 1':
      2,18,88,635    cpu_atom/cycles/          (96.6
      2,80,08,188    cpu_core/cycles/         (3.37
      2,17,99,872    cpu_atom/instructions/   #  1.00  linsn per cycle (96.6
      2,53,56,126    cpu_core/instructions/   #  0.91  linsn per cycle (3.37
      56,591         cpu_atom/cache-misses/   (96.6
      1,73,887       cpu_core/cache-misses/   (3.37
<not supported>  cpu_atom/L1-dcache-load-misses/
      1,77,421       cpu_core/L1-dcache-load-misses/ (3.37
      2             context-switches          (3.37
121.036659255 seconds time elapsed
 0.000000000 seconds user
 0.006726000 seconds sys

iiitd@iiitd-ThinkCentre-M70s-Gen-3: ~/Desktop/GR5_PA02$ ./MT25067_PartA1_Client 4096 5000
*** Part A1: Two-Copy Client ***
Expecting 5000 messages of 4096 bytes
Connected to server at 127.0.0.1:8080
Received 1000 messages...
Received 2000 messages...
Received 3000 messages...
Received 4000 messages...
Received 5000 messages...
*** Results ***
Messages received: 5000
Total bytes: 20480000
Time elapsed: 0.006 sec
Throughput: 28248.28 Mbps
Average latency: 1.16 µs
iiitd@iiitd-ThinkCentre-M70s-Gen-3: ~/Desktop/GR5_PA02$ 

```

Figure 1: Perf statistics showing CPU cycles, cache misses, and context switches for A1 implementation

Concurrent Client Handling:

```
iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/GRS_PA02$ ./MT25067_PartA1_Server 4096 5000 4
== Part A1: Two-Copy Server ==
Message size: 4096 bytes
Messages per client: 5000
Max clients: 4
Server listening on port 8080...
Client 1 connected
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (36984.20 Mbps)
Client 2 connected
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (41700.18 Mbps)
Client 3 connected
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (43481.95 Mbps)
Client 4 connected
All 4 clients accepted. Waiting for transfers to complete...
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (44401.08 Mbps)

== Final Statistics ==
Total bytes sent: 81920000
Total time: 0.016 sec
Average throughput: 41433.90 Mbps
iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/GRS_PA02$
```

Figure 2: Server successfully accepting and processing 4 concurrent clients

Automated Experiment Execution:

```
iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/GRS_PA02$ bash MT25067_PartC_AutomationScript.sh
Feb 6 22:35

[INFO] =====
[INFO] MT25067 Automated Experiment Runner
[INFO] =====

[INFO] Compiling all implementations...
[DEBUG] Cleaned previous build.
gcc -Wall -fextra -O2 -pthread -o MT25067_PartA1_Server MT25067_PartA1_Server.c -pthread
gcc -Wall -fextra -O2 -pthread -o MT25067_PartA1_Client MT25067_PartA1_Client.c -pthread
gcc -Wall -fextra -O2 -pthread -o MT25067_PartA2_Server MT25067_PartA2_Server.c -pthread
gcc -Wall -fextra -O2 -pthread -o MT25067_PartA2_Client MT25067_PartA2_Client.c -pthread
gcc -Wall -fextra -O2 -pthread -o MT25067_PartA3_Server MT25067_PartA3_Server.c -pthread
gcc -Wall -fextra -O2 -pthread -o MT25067_PartA3_Client MT25067_PartA3_Client.c -pthread
Build complete. All parts compiled successfully.
[INFO] Compilation successful
[INFO] CSV file: MT25067_ExperimentData.csv (main directory)
[INFO] Temp directory: experiment_results

[INFO] Configuration:
[INFO]   Message sizes: 256 1024 4096 16384 bytes
[INFO]   Thread counts: 1 2 4 8
[INFO]   Messages per client: 5000 (FIXED: increased for stability)
[INFO]   Total experiments: 48

[WARN] NOTE: This will take approximately 24 minutes
[WARN] Press Ctrl+C within 5 seconds to cancel...

[FZ] [INFO]
[INFO] Progress: 1 / 48
[INFO] _____
[INFO] Running experiment: A1_256B_1T
[INFO] Starting server for A1_256B_1T on port 8080
[DEBUG] Server PID: 89505
[INFO] Server ready, starting 1 client(s)
[DEBUG] Started 1 clients
[INFO] Clients completed, waiting for server to finish...
[INFO] ✓ Experiment A1_256B_1T completed successfully

[?] [INFO]
[INFO] Progress: 2 / 48
[INFO] _____
[INFO] Running experiment: A1_256B_2T
[INFO] Starting server for A1_256B_2T on port 8080
[DEBUG] Server PID: 89573
[INFO] Server ready, starting 2 client(s)
[DEBUG] Started 2 clients
[INFO] Clients completed, waiting for server to finish...
[INFO] ✓ Experiment A1_256B_2T completed successfully

[?] [INFO]
[INFO] Progress: 3 / 48
[INFO] _____
[INFO] Running experiment: A1_256B_4T
[INFO] Starting server for A1_256B_4T on port 8080
```

Figure 3: Automation script running all 48 experiment combinations

Part A3: Zero-Copy with MSG_ZEROCOPY

Kernel Behavior:

```
sendmsg(fd, msg, MSG_ZEROCOPY)
↓
```

```

Kernel pins user pages (get_user_pages)
    ↓
Marks pages as "in-flight" (reference counting)
    ↓
[On physical NIC]: DMA engine reads directly from pinned pages
[On loopback]:     Kernel detects loopback → falls back to memcpy
    ↓
Async completion → notification via MSG_ERRQUEUE

```

Why it's slower on localhost: No actual DMA occurs; page pinning overhead + completion polling overhead > memcpy cost.

4. Experimental Results

4.1 Measurement Parameters

Parameter	Values
Message Sizes	256B, 1KB, 4KB, 16KB
Thread Counts	1, 2, 4, 8
Messages per Client	5000
Total Experiments	48 (3 implementations × 4 sizes × 4 threads)

4.2 Key Findings Summary

Metric	A1 (Best)	A2 (Best)	A3 (Best)
Peak Throughput	57.79 Gbps	56.69 Gbps	31.81 Gbps
Lowest Latency	0.42 µs	0.67 µs	2.72 µs
Lowest LLC Misses	83,877	132,763 (at 4KB)	161,703
Optimal Thread Count	2	4	1
Best Message Size	16KB	16KB	16KB


Winner by Category:

- **Small messages (<1KB):** A1 (Two-Copy) — lowest syscall overhead
 - **Medium messages (4-16KB):** A1/A2 tied — both perform well
 - **Large messages (>64KB):** A2 (One-Copy) likely wins; A3 would win with real NIC
-

5. Plots and Visualizations

All plots generated using `matplotlib` with hardcoded experimental data (as per assignment requirements).


Plot 1: Throughput vs Message Size



Key Observations:

- Throughput scales with message size for A1 and A2
 - A3 severely underperforms due to localhost overhead
 - 2-thread configuration gives best balance for most message sizes
-


Plot 2: Latency vs Thread Count



Key Observations:

- Latency increases sharply at 8 threads (context switching)
- A3 has consistently higher latency (3-6μs) due to completion polling
- Optimal latency achieved with 1-2 threads for all implementations


Plot 3: Cache Misses vs Message Size



Key Observations:

- LLC misses: A2 shows 65% reduction at 4KB (one-copy benefit)
 - L1 misses: More variable due to instruction cache interference
 - Larger messages → more cache pressure → more misses
-


Plot 4: CPU Cycles per Byte



Key Observations:

- A3 uses 3-5x more cycles per byte (page pinning overhead)
 - A1 and A2 converge at larger message sizes (~15-20 cycles/byte)
 - Efficiency improves with message size (fixed syscall cost amortized)
-

Plot 5: Overall Comparison (16KB messages)



Key Observations:

- 4 threads is the "sweet spot" for A1/A2 before context switch overhead
- A3 throughput remains poor across all thread counts (localhost limitation)
- Cache misses explode at 8 threads (3.5M for A1 vs 850K at 4 threads)

6. Code Quality & Structure

Modularity

- Each implementation in separate files (A1, A2, A3)
- Clear separation of client and server logic
- Reusable message structure across all implementations

Comments & Documentation

- Every function has purpose comments
- Critical sections marked (e.g., `// TWO-COPY:`, `// ZERO-COPY:`)
- Inline explanations for non-obvious operations (page pinning, iovec setup)

Error Handling

- All syscalls checked for errors (`if (ret < 0) perror(...)`)
- Graceful cleanup with `goto` labels

- Mutex lock/unlock safety (no deadlocks)

Alignment & Style

- Consistent 4-space indentation
 - Descriptive variable names (`bytes_sent_total`, not `bst`)
 - Logical grouping of related code blocks
-

7. AI Usage Declaration

Components Where AI Was Used

1. Bash Automation Script

Tool: Gemini Flash 1.5

Prompts Used:

1. "Write a bash script to compile C server/client programs, run them with different message sizes and thread counts, and collect perf stat output including cycles and cache-misses into a CSV file."
2. "Add error handling to check if ports are in use before starting servers, and kill any existing processes on those ports."
3. "Fix the script to handle perf output parsing when numbers have commas (e.g., 1,234,567 cycles)."

What Was Generated:

- Loop structure for iterating over experimental parameters
- `lsof` and `kill` commands for port cleanup
- CSV parsing logic with `awk` and `grep`
- `tr -d ','` fix for comma-separated numbers in perf output

What I Modified:

- Increased `NUM_MESSAGES` from 1000 to 5000 for stable profiling data
- Added validation function `validate_perf_output()` to catch incomplete data
- Extended server wait time from 2s to 5s for proper synchronization
- Added progress bar and colored output for better UX

2. Python Plotting Script

Tool: Gemini Flash 1.5 / ChatGPT-4

Prompts Used:

1. "Create a Python script using matplotlib to plot Throughput vs Message Size from a CSV with columns: Implementation, MessageSize, Throughput."
2. "Generate 4 subplots showing different thread counts, with log scale on x-axis for message sizes."
3. "Add a footer to each plot with system configuration details."

What Was Generated:

- Basic matplotlib plotting structure

- Subplot layout (2x2 grid)
- Legend and axis labeling

What I Modified:

- Hardcoded experimental data directly in Python (per assignment requirement)
- Converted Mbps to Gbps (divided by 1000)
- Fixed cycle data to use millions (e.g., 13.445259 instead of 13445259)
- Added 5th plot for comprehensive comparison
- Customized colors, markers, and annotations

3. Analysis Assistance (Part E)

Tool: Claude 3.5 Sonnet

Prompts Used:

1. "Explain why MSG_ZEROCOPY might perform worse than regular send() on a localhost loopback interface."
2. "What are the micro-architectural reasons for context switch overhead affecting network throughput?"

What AI Suggested:

- Page pinning overhead explanation
- Loopback fallback behavior (no real DMA on `lo` interface)
- TLB flushing and cache warmup cost per context switch

How I Used It:

- Verified suggestions against Linux kernel documentation ([Documentation/networking/msg_zerocopy.rst](#))
- Cross-referenced with course notes on cache coherency protocols
- Used as starting point, then rewrote in my own words with specific experimental evidence

Components NOT Generated by AI

- **All C code** (A1, A2, A3 server/client implementations) — 100% hand-written
- **Makefile** — written manually
- **Experimental design** — decided message sizes, thread counts based on assignment requirements
- **Analysis insights** — AI provided background theory, but all data interpretation and conclusions are mine
- **Report structure** — organized based on assignment rubric

Honesty Declaration

I understand that failure to properly declare AI usage constitutes plagiarism. The above declaration is complete and accurate to the best of my knowledge. I am prepared to explain every line of code and every analysis conclusion during the viva.

8. GitHub Repository

URL: https://github.com/dewansh3255/GRS_PA02

Visibility: Public

Folder Name: GRS_PA02 (as required)

Repository Contents:

```
GRS_PA02/
└── MT25067_PartA1_Server.c
└── MT25067_PartA1_Client.c
└── MT25067_PartA2_Server.c
└── MT25067_PartA2_Client.c
└── MT25067_PartA3_Server.c
└── MT25067_PartA3_Client.c
└── MT25067_PartC_AutomationScript.sh
└── MT25067_PartD_Plots.py
└── MT25067_PartE_Analysis.md
└── MT25067_ExperimentData.csv
└── MT25067_Report.md (this file)
└── Makefile
└── README.md
```

Note: No binary files, no PNG plots (plots are embedded in this report only), no subfolders.

9. Viva Preparation Notes

Questions I Expect

Q1: Explain the difference between A1, A2, and A3 at the kernel level.

A: A1 uses `send()` which copies from user buffer → kernel buffer. A2 uses `sendmsg()` with `iovec`, allowing kernel to gather from scattered user buffers in one copy. A3 uses `MSG_ZEROCOPY` which pins user pages and attempts DMA (fails on loopback, falls back to copy).

Q2: Why is context switching so expensive?

A: Each switch requires saving/restoring registers (~100 cycles), flushing TLB (~500 cycles), and warming up caches (~5000 cycles). The real cost is the cache warmup phase slowing down subsequent operations.

Q3: What would change if you ran this on a real network (not localhost)?

A: A3 would perform much better because the NIC's DMA engine could read directly from pinned user pages, avoiding CPU copy. Zero-copy would win for large messages (>64KB).

Q4: How did you ensure thread safety in your server?

A: Used a global statistics structure with a `pthread_mutex_t` lock. Each thread acquires the lock before updating shared counters (`total_bytes_sent`), then releases it.

Q5: What's the bottleneck at 8 threads?

A: The i7-12700 has 8 P-cores + 4 E-cores. At 8 threads, we saturate the P-cores and compete for resources, causing context switching (122/sec vs 14/sec at 4 threads). The mutex lock also becomes a serialization point.

10. Deliverables Checklist (Final)

- **MT25067_PA02.zip** uploaded to Google Classroom
 - **GRS_PA02/** folder pushed to GitHub (public repo)
 - All files follow naming convention **MT25067_PartX_...**
 - Makefile has roll number comment at top
 - README has roll number comment at top
 - No binary files in submission
 - No PNG plots in zip (only in report)
 - CSV file present with raw data
 - Python script has hardcoded data (no CSV reading)
 - AI usage declared in detail (tools, prompts, modifications)
 - Report includes GitHub URL
-

11. Conclusion

This assignment provided hands-on experience with:

- **Low-level network I/O primitives** (send, sendmsg, MSG_ZEROCOPY)
- **Performance profiling** using Linux perf
- **Cache behavior analysis** (LLC vs L1 misses)
- **Multithreaded programming** challenges (context switching, lock contention)
- **System optimization** trade-offs (zero-copy isn't always better)

Key Takeaway: The "best" implementation depends on workload characteristics. Small messages favor simple primitives (A1), medium messages benefit from reduced copies (A2), and zero-copy (A3) requires specific conditions (large messages + hardware offload) to shine.

I am confident in my understanding of all implementations and analysis, and ready to defend this work during the viva.

End of Report

Submitted by: MT25067 (Dewansh Khandelwal)

Date: February 6, 2026