

# Analysis of Network I/O Primitives Using Perf

---

**Student:** Dewansh Khandelwal (MT25067)

**Course:** CSE638 - Graduate Systems

**Date:** February 6, 2026

---

## Abstract

This report presents a comprehensive experimental analysis of three network I/O paradigms: two-copy (baseline), one-copy (scatter-gather), and zero-copy (MSG\_ZEROCOPY) socket communication. Using the Linux `perf` tool, we profiled CPU cycles, cache behavior, and context switches across varying message sizes (256B-16KB) and thread counts (1-8). Our findings reveal that zero-copy underperforms on localhost due to page pinning overhead without DMA benefits, one-copy reduces LLC misses by 65% through elimination of serialization buffers, and optimal thread count is bounded by cache coherency and context switching costs. The crossover point for one-copy vs two-copy occurs at 8-16KB message sizes.

---

## 1. Introduction

### 1.1 Motivation

Network I/O operations traditionally involve multiple data copies between user space and kernel space, consuming CPU cycles and memory bandwidth. Modern Linux provides optimizations like scatter-gather I/O (`sendmsg` with `iovec`) and zero-copy (`MSG_ZEROCOPY`) to reduce this overhead. Understanding when each approach provides actual performance benefits requires empirical analysis under controlled conditions.

### 1.2 Experimental Setup

#### Hardware Configuration:

- **CPU:** Intel Core i7-12700 (12th Gen)
  - 8 Performance cores (P-cores) @ 3.6-4.9 GHz
  - 4 Efficiency cores (E-cores) @ 2.7-3.6 GHz
  - 25MB shared L3 cache
- **RAM:** DDR4 (sufficient for workload)
- **Network:** Localhost loopback interface (`lo`)

#### Software Environment:

- **OS:** Ubuntu 22.04.5 LTS
- **Kernel:** 6.8.0-90-generic
- **Compiler:** GCC 11.4.0 with `-O2 -pthread`
- **Profiler:** Linux perf (perf\_event subsystem)


#### Experimental Parameters:

Parameter	Values
Message Sizes	256B, 1KB, 4KB, 16KB
Thread Counts	1, 2, 4, 8
Messages per Client	5000
Total Experiments	48 (3 implementations × 4 sizes × 4 threads)

## 2. Implementation Details

### 2.1 Two-Copy Implementation (A1 - Baseline)

#### Architecture:



#### Key Code Snippet:


```

// Copy #1: User-space serialization
int serialize_message(Message *msg, char *buffer, int buffer_size) {
    int offset = 0;
    int field_size = buffer_size / 8;
    memcpy(buffer + offset, msg->field1, field_size); offset += field_size;
    memcpy(buffer + offset, msg->field2, field_size); offset += field_size;
    // ... repeat for all 8 fields
    return offset;
}

// Copy #2: Kernel-space copy via send()
send(client_fd, send_buffer, message_size, 0);
    
```

### 2.2 One-Copy Implementation (A2 - Scatter-Gather)

#### Architecture:


**Key Code Snippet:**

```

// Setup iovec to point directly to scattered buffers
struct iovec iov[8];
iov[0].iov_base = msg->field1; // No memcpy!
iov[0].iov_len = field_size;
// ... setup all 8 iovecs

// Kernel performs scatter-gather I/O
struct msghdr msghdr;
msghdr.msg_iov = iov;
msghdr.msg_iovlen = 8;
sendmsg(client_fd, &msghdr, 0);
    
```

## 2.3 Zero-Copy Implementation (A3 - MSG\_ZEROCOPY)

**Architecture:****Key Code Snippet:**

```

// Enable zero-copy on socket
int optval = 1;
    
```

```

setsockopt(client_fd, SOL_SOCKET, SO_ZEROCOPY, &optval, sizeof(optval));

// Send with MSG_ZEROCOPY flag
sendmsg(client_fd, &msghdr, MSG_ZEROCOPY);

// Must poll for completion
while (pending_completions > 0) {
    recvmsg(sock_fd, &msg, MSG_ERRQUEUE | MSG_DONTWAIT);
}

```

## 3. Part B: Profiling and Measurement Methodology

### 3.1 Measurement Infrastructure

All experiments were profiled using Linux `perf stat` with the following event counters:

```

sudo perf stat -e cycles,instructions,cache-misses,L1-dcache-load-
misses,context-switches \
./MT25067_PartA1_Server

```

#### Perf Events Collected:

- **cycles**: Total CPU cycles consumed
- **instructions**: Total instructions executed (for IPC calculation)
- **cache-misses**: Last Level Cache (LLC/L3) misses
- **L1-dcache-load-misses**: L1 data cache load misses
- **context-switches**: Thread context switches

### 3.2 Application-Level Metrics

Throughput and latency calculated in client code using `gettimeofday()`:

```

struct timeval start, end;
gettimeofday(&start, NULL);
// ... receive data ...
gettimeofday(&end, NULL);

double elapsed = (end.tv_sec - start.tv_sec) +
                 (end.tv_usec - start.tv_usec) / 1e6;

double throughput_gbps = (total_bytes * 8.0) / (elapsed * 1e9);
double avg_latency_us = (elapsed * 1e6) / messages_received;

```

### 3.3 Data Collection Summary

**Total measurements:** 48 experiments

- **Implementations:** 3 (A1, A2, A3)
- **Message sizes:** 4 (256B, 1KB, 4KB, 16KB)
- **Thread counts:** 4 (1, 2, 4, 8)
- **Repetitions:** Each experiment run once with 5000 messages per client

**Sample Collected Data (A1, 4KB, 4 threads):**

Metric	Value	Unit
Throughput	42.26	Gbps
Latency	0.78	µs
Total Bytes	20,480,000	bytes
CPU Cycles	150,723,683	cycles
Instructions	185,716,319	instructions
LLC Misses	337,977	misses
L1 Misses	805,432	misses
Context Switches	5	switches
Time Elapsed	6.03	seconds

All raw measurements stored in [MT25067\\_ExperimentData.csv](#) for reproducibility.

---

## 4. Experimental Results

### 4.1 Throughput Analysis

**Throughput vs Message Size (4 threads):**

Message Size	A1 (Gbps)	A2 (Gbps)	A3 (Gbps)	Winner
256B	2.85	3.04	0.51	A2
1KB	15.94	12.27	1.90	A1
4KB	42.26	31.12	0.75	A1
16KB	53.76	56.69	26.50	<b>A2</b>

**Key Finding:** A2 (one-copy) achieves highest throughput at 16KB messages (56.69 Gbps), while A3 (zero-copy) performs 5-10x worse across all sizes.

---

### 4.2 Cache Behavior Analysis

**Cache Misses for 4KB Messages, 1 Thread:**

Metric	A1 (Two-Copy)	A2 (One-Copy)	A3 (Zero-Copy)	Best
LLC Misses	385,006	132,763 (-65%)	182,509	A2 ✓
L1 Misses	133,815	250,227 (+87%)	569,753	A1
CPU Cycles	47.3M	55.4M	249.5M	A1

### Analysis:

The 65% reduction in LLC (L3) misses for A2 is explained by:

#### 1. Memory Footprint Reduction:

- A1: 8 fields (4KB) + serialization buffer (4KB) = **8KB working set**
- A2: 8 fields (4KB) only = **4KB working set**

#### 2. Cache Line Reuse:

- A1 serialization causes cache eviction
- A2 kernel reads directly from original locations

## 4.3 Thread Scaling and Context Switching

### Context Switching Impact (16KB messages):

Threads	Throughput (Gbps)	Context Switches	LLC Misses	Analysis
1	55.01	3	227K	Baseline
2	57.79	6	641K	Optimal ✓
4	53.76	14	849K	Diminishing
8	14.83 (-73%)	122 (+771%)	3.5M	<b>Collapse</b>

**Critical Finding:** The 73% throughput drop from 4→8 threads is caused by context switch storm and cache thrashing.

## 4.4 CPU Efficiency

### CPU Cycles per Byte Transferred (4 threads):

Message Size	A1 (cycles/byte)	A2 (cycles/byte)	A3 (cycles/byte)
256B	63.04	73.16	582.04
1KB	13.60	22.04	187.15
4KB	7.36	9.07	54.43
16KB	5.37	5.03	14.78


**Insight:** As message size increases, per-byte cost converges for A1/A2 (~5 cycles/byte), while A3 remains 3x worse due to page pinning overhead.

---

## 5. Experimental Evidence

### 5.1 Perf Profiling Output

#### Sample Output (4KB, 1 thread):



```

Activities Terminal Feb 6 22:30
iiitd@iiitd-ThinkCentre-M70s-Gen-3: ~/Desktop/GRS_PA02$ sudo perf stat -e cycles,instructions,cache-misses,L1-dcache-load-misses,context-switches ./MT25067_PartA1_Server 4096 5000 1
*** Part A1: Two-Copy Server ===
Message size: 4096 bytes
Messages per client: 5000
Max clients: 1
Server listening on port 8080...
Clients connected...
All clients accepted. Waiting for transfers to complete...
[Thread 129918281184832] Handling Client, sending 5000 messages of 4096 bytes
[Thread 129918281184832] Sent 20480000 bytes in 0.006 sec (29356.75 Mbps)

*** Final Statistics ===
Total bytes sent: 20480000
Total time: 0.006 sec
Average throughput: 29356.75 Mbps

Performance counter stats for './MT25067_PartA1_Server 4096 5000 1':
      2,18,88,635    cpu_atom/cycles/          (96.6
      2,80,08,188    cpu_core/cycles/         (3.37
      2,17,99,872   cpu_atom/instructions/ # 1.00  insn per cycle (96.6
      2,53,56,126   cpu_core/instructions/ # 0.91  insn per cycle (3.37
      56,591        cpu_atom/cache-misses/     (96.6
      1,73,887      cpu_core/cache-misses/    (3.37
<not supported>  cpu_atom/L1-dcache-load-misses/
      1,77,421      cpu_core/L1-dcache-load-misses/ (3.37
      2             context-switches          (96.6

121.036659255 seconds time elapsed
0.000000000 seconds user
0.006726000 seconds sys

iiitd@iiitd-ThinkCentre-M70s-Gen-3: ~/Desktop/GRS_PA02$ 

iiitd@iiitd-ThinkCentre-M70s-Gen-3: ~/Desktop/GRS_PA02$ ./MT25067_PartA1_Client 4096 5000
*** Part A1: Two-Copy Client ===
Expecting 5000 messages of 4096 bytes
Connected to server at 127.0.0.1:8080
Received 1000 messages...
Received 2000 messages...
Received 3000 messages...
Received 4000 messages...
Received 5000 messages...

*** Results ***
Messages received: 5000
Total bytes: 20480000
Time elapsed: 0.006 sec
Throughput: 29356.75 Mbps
Average latency: 1.16 µs
iiitd@iiitd-ThinkCentre-M70s-Gen-3: ~/Desktop/GRS_PA02$ 

```

Figure 1: Perf statistics showing 2 context switches and LLC/L1 cache misses for single-threaded execution

---

### 5.2 Multi-Client Server Execution


```
iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/GRS_PA02$ ./MT25067_PartA1_Server 4096 5000 4
== Part A1: Two-Copy Server ==
Message size: 4096 bytes
Messages per client: 5000
Max clients: 4
Server listening on port 8080...
Client 1 connected
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (36984.20 Mbps)
Client 2 connected
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (41700.18 Mbps)
Client 3 connected
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (43481.95 Mbps)
Client 4 connected
All 4 clients accepted. Waiting for transfers to complete...
[Thread 131431959361088] Handling client, sending 5000 messages of 4096 bytes
[Thread 131431959361088] Sent 20480000 bytes in 0.004 sec (44401.08 Mbps)

== Final Statistics ==
Total bytes sent: 81920000
Total time: 0.016 sec
Average throughput: 41433.90 Mbps
iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/GRS_PA02$
```

Figure 2: Server successfully handling 4 concurrent clients with separate threads

**Key Observation:** Each client handled by independent thread, achieving 41.4 Gbps aggregate throughput.

### 5.3 Automation Script Execution



```
iiitd@iiitd-ThinkCentre-M70s-Gen-3:~/Desktop/GRS_PA02$ bash MT25067_PartC_AutomationScript.sh
[INFO] =====
[INFO] MT25067 Automated Experiment Runner
[INFO] =====
[INFO] Compiling all implementations...
[DEBUG] Cleaned previous builds.
gcc -Wall -Wextra -O2 -pthread -o MT25067_PartA1_Server MT25067_PartA1_Server.c -pthread
gcc -Wall -Wextra -O2 -pthread -o MT25067_PartA1_Client MT25067_PartA1_Client.c -pthread
gcc -Wall -Wextra -O2 -pthread -o MT25067_PartA2_Server MT25067_PartA2_Server.c -pthread
gcc -Wall -Wextra -O2 -pthread -o MT25067_PartA2_Client MT25067_PartA2_Client.c -pthread
gcc -Wall -Wextra -O2 -pthread -o MT25067_PartA3_Server MT25067_PartA3_Server.c -pthread
gcc -Wall -Wextra -O2 -pthread -o MT25067_PartA3_Client MT25067_PartA3_Client.c -pthread
Build completed! All packages build successfully.
[INFO] Compilation successful
[INFO] CSV file: MT25067_ExperimentData.csv (main directory)
[INFO] Temp directory: experiment_results
[INFO] Configuration:
[INFO]   Message sizes: 256 1024 4096 16384 bytes
[INFO]   Thread counts: 1 2 4 8
[INFO]   Messages per client: 5000 (FIXED: Increased for stability)
[INFO]   Total experiments: 48
[WARN] NOTE: This will take approximately 24 minutes
[WARN] Press Ctrl+C within 5 seconds to cancel...
[INFO] Progress: 1 / 48
[INFO] Running experiment: A1_256B_1T
[INFO] Starting server for A1_256B_1T on port 8080
[DEBUG] Server PID: 89505
[INFO] Server ready, starting 1 client(s)
[DEBUG] Started 1 clients
[INFO] Clients completed, waiting for server to finish...
[INFO] ✓ Experiment A1_256B_1T completed successfully
[INFO] Progress: 2 / 48
[INFO] Running experiment: A1_256B_2T
[INFO] Starting server for A1_256B_2T on port 8080
[DEBUG] Server PID: 89573
[INFO] Server ready, starting 2 client(s)
[DEBUG] Started 2 clients
[INFO] Clients completed, waiting for server to finish...
[INFO] ✓ Experiment A1_256B_2T completed successfully
[INFO] Progress: 3 / 48
[INFO] Running experiment: A1_256B_4T
[INFO] Starting server for A1_256B_4T on port 8080
```

Figure 3: Automated experiment runner executing all 48 test combinations

## 6. Detailed Analysis

## 6.1 Why Zero-Copy Fails on Localhost

### Overhead Breakdown (4KB message):

#### Zero-Copy Send:

1. get_user_pages():	~2000 cycles (TLB walk)
2. Reference counting:	~1000 cycles (atomic ops)
3. memcpy (loopback):	~1500 cycles (fallback)
4. Completion notification:	~2000 cycles (polling)
Total:	~6500 cycles

#### Two-Copy Send:

1. User memcpy:	~1500 cycles
2. send() syscall:	~1000 cycles
3. Kernel copy_from_user():	~1500 cycles
Total:	~4000 cycles

Penalty:  $6500 / 4000 = 1.625 \times$  slower

### When Zero-Copy Wins:

- Physical 10GbE/40GbE NIC with DMA
- Large messages (>64KB)
- High-throughput workloads

## 6.2 One-Copy vs Two-Copy Crossover

### Crossover Analysis:

Message Size	A1 Advantage	A2 Advantage	Reason
256B	✓ +7.7%	-	sendmsg overhead > memcpy cost
1KB	✓ +30%	-	Syscall setup dominates
4KB	✓ +6%	-	Near crossover
<b>16KB</b>	-	✓ +5.5%	<b>Memcpy cost dominates</b>

### Mathematical Model:

```
Crossover when:
memcpy_cost(size) = iovec_setup_cost + overhead_diff
0.375 × size = 300 + kernel_overhead
size ≈ 10–16KB
```

## 6.3 8-Thread Performance Collapse

## Root Cause Analysis:

### 1. Hardware Asymmetry:

```
i7-12700: 8 P-cores (fast) + 4 E-cores (slow)  
4 threads: All fit on P-cores → optimal  
8 threads: Spill to E-cores → 30% slower + migration overhead
```

### 2. Context Switch Storm:

```
4 threads: 14 switches/sec  
8 threads: 122 switches/sec (+771%)  
  
Per-switch cost: ~5600 cycles  
122 × 5600 = 683K cycles/sec wasted
```


### 3. Cache Thrashing:

```
LLC Misses:  
4 threads: 849K  
8 threads: 3.5M (+312%)  
  
MESI protocol overhead: Cache line ping-pong between cores
```

---


## 7. Visualization of Results

Plot 1: Throughput vs Message Size




Shows A1 and A2 scaling with message size, A3 consistently underperforming

## Plot 2: Latency vs Thread Count




*Demonstrates latency spike at 8 threads due to context switching*

### Plot 3: Cache Misses vs Message Size




*Highlights 65% LLC miss reduction in A2 vs A1*

### Plot 4: CPU Cycles per Byte



Shows efficiency convergence at large message sizes

### Plot 5: Overall Comparison (16KB)



Comprehensive view of all metrics at optimal message size

## 8. Conclusions

### 8.1 Key Findings

- Zero-Copy Paradox:** MSG\_ZEROCOPY on localhost has **negative benefit** (2-10x slower)
- Cache Optimization:** One-copy reduces LLC misses by 65% despite increasing L1 misses
- Thread Count:** Optimal at 2-4 threads; 8 threads causes 73% throughput collapse
- Message Size Crossover:** One-copy wins at 8-16KB boundary

### 8.2 Practical Recommendations

Scenario	Recommended Implementation
Small messages (<1KB)	A1 (send/recv)
Medium messages (1-16KB)	A2 (sendmsg + iovec)
Large messages + real NIC	A3 (MSG_ZEROCOPY)
Thread count	≤ P-core count (8 for i7-12700)

## 9. AI Usage Declaration

### Components Where AI Was Used

#### 1. Bash Automation Script

**Tool:** Gemini Pro 3

**Prompts Used:**

1. "Write a bash script to compile C server/client programs, run them with different message sizes and thread counts, and collect perf stat output including cycles and cache-misses into a CSV file."
2. "Add error handling to check if ports are in use before starting servers, and kill any existing processes on those ports."
3. "Fix the script to handle perf output parsing when numbers have commas (e.g., 1,234,567 cycles)."

**What Was Generated:**

- Loop structure for iterating over experimental parameters
- `lsof` and `kill` commands for port cleanup
- CSV parsing logic with `awk` and `grep`
- `tr -d ','` fix for comma-separated numbers in perf output

**What I Modified:**

- Increased `NUM_MESSAGES` from 1000 to 5000 for stable profiling data
- Added validation function `validate_perf_output()` to catch incomplete data
- Extended server wait time from 2s to 5s for proper synchronization
- Added progress bar and colored output for better UX

#### 2. Python Plotting Script

**Tool:** Gemini Flash / ChatGPT-4

**Prompts Used:**

1. "Create a Python script using matplotlib to plot Throughput vs Message Size from a CSV with columns: Implementation, MessageSize, Throughput."
2. "Generate 4 subplots showing different thread counts, with log scale on x-axis for message sizes."
3. "Add a footer to each plot with system configuration details."

**What Was Generated:**

- Basic matplotlib plotting structure
- Subplot layout (2x2 grid)
- Legend and axis labeling

**What I Modified:**

- Hardcoded experimental data directly in Python (per assignment requirement)
- Converted Mbps to Gbps (divided by 1000)
- Fixed cycle data to use millions (e.g., 13.445259 instead of 13445259)
- Added 5th plot for comprehensive comparison

- Customized colors, markers, and annotations

### 3. Analysis Assistance (Part E)

**Tool:** Claude 3.5 Sonnet

**Prompts Used:**

1. "Explain why MSG\_ZEROCOPY might perform worse than regular send() on a localhost loopback interface."
2. "What are the micro-architectural reasons for context switch overhead affecting network throughput?"

**What AI Suggested:**

- Page pinning overhead explanation
- Loopback fallback behavior (no real DMA on `lo` interface)
- TLB flushing and cache warmup cost per context switch

**How I Used It:**

- Verified suggestions against Linux kernel documentation ([Documentation/networking/msg\\_zerocopy.rst](#))
- Cross-referenced with course notes on cache coherency protocols
- Used as starting point, then rewrote in my own words with specific experimental evidence

Components NOT Generated by AI

- **All C code** (A1, A2, A3 server/client implementations) — 100% hand-written
- **Experimental design** — decided message sizes, thread counts based on assignment requirements
- **Analysis insights** — AI provided background theory, but all data interpretation and conclusions are mine
- **Report structure** — organized based on assignment rubric

---

## 10. GitHub Repository

**URL:** [https://github.com/dewansh3255/GRS\\_PA02](https://github.com/dewansh3255/GRS_PA02)

**Visibility:** Public

**Folder Name:** `GRS_PA02` (as required)

**Repository Contents:**

```
GRS_PA02/
├── MT25067_PartA1_Server.c
├── MT25067_PartA1_Client.c
├── MT25067_PartA2_Server.c
├── MT25067_PartA2_Client.c
├── MT25067_PartA3_Server.c
├── MT25067_PartA3_Client.c
└── MT25067_PartC_AutomationScript.sh
```

```
└── MT25067_PartD_Plots.py  
└── MT25067_PartE_Analysis.md  
└── MT25067_ExperimentData.csv  
└── MT25067_Report.md (this file)  
└── Makefile  
└── README.md
```

**Note:** No binary files, no PNG plots (plots are embedded in this report only), no subfolders.

---

## 11. Conclusion

This assignment provided hands-on experience with:

- **Low-level network I/O primitives** (send, sendmsg, MSG\_ZEROCOPY)
- **Performance profiling** using Linux perf
- **Cache behavior analysis** (LLC vs L1 misses)
- **Multithreaded programming** challenges (context switching, lock contention)
- **System optimization** trade-offs (zero-copy isn't always better)

**Key Takeaway:** The "best" implementation depends on workload characteristics. Small messages favor simple primitives (A1), medium messages benefit from reduced copies (A2), and zero-copy (A3) requires specific conditions (large messages + hardware offload) to shine.

I am confident in my understanding of all implementations and analysis, and ready to defend this work during the viva.

---

## End of Report

*Submitted by: MT25067 (Dewansh Khandelwal)*

*Date: February 7, 2026*